

PDF Vectorization and Similarity Search Demo

This sample document demonstrates a PDF processing pipeline.

The pipeline loads PDF pages, splits text into overlapping chunks, creates vector embeddings, and stores them in a ChromaDB database.

FastAPI provides a query endpoint that returns relevant text chunks, page numbers, and vector search scores.

The project uses LangChain and LangGraph and does not use an LLM.